

Extending Compiler Support for the BrainScaleS Plasticity Processor

Bachelor's Thesis Presentation

Arthur Heimbrecht

March 14, 2017

Contents

PPU Architecture

GCC Structure

Extending GCC for the PPU

Results

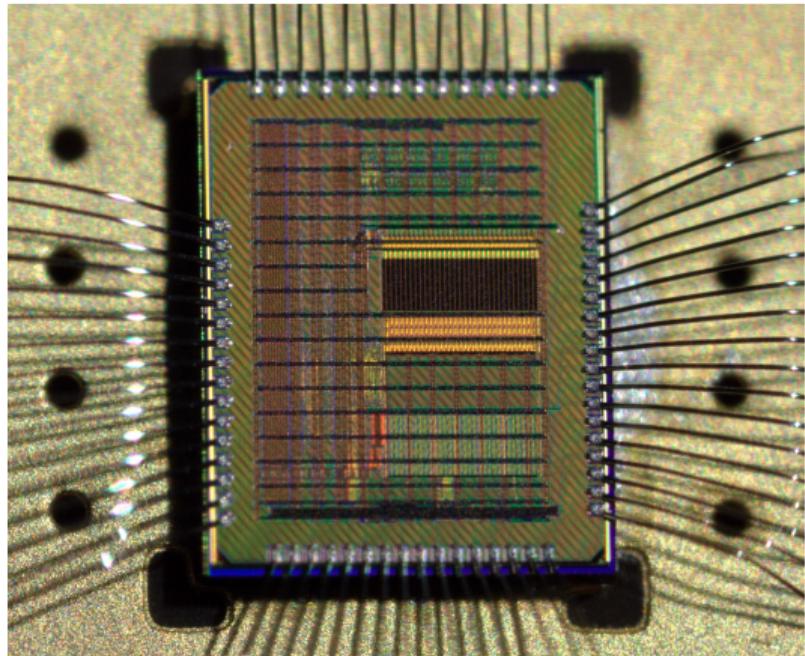


Figure 1: Photograph of HICANN-DLS chip, Friedmann et al.

HICANN-DLS

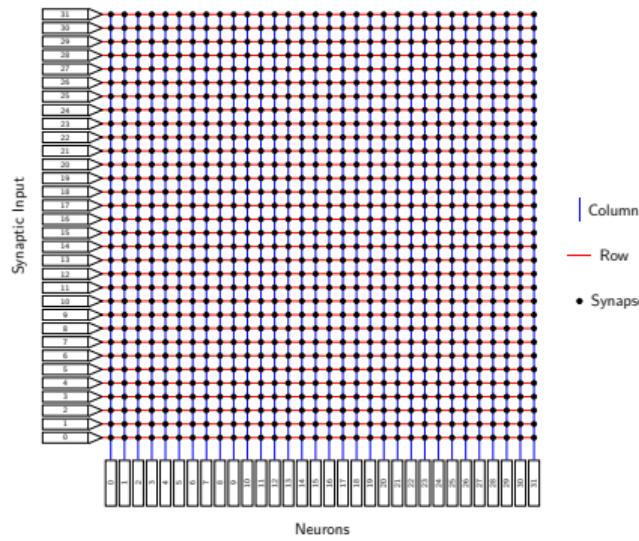


Figure 2: Schematic Representation of the Synapse Array on HICANN-DLS

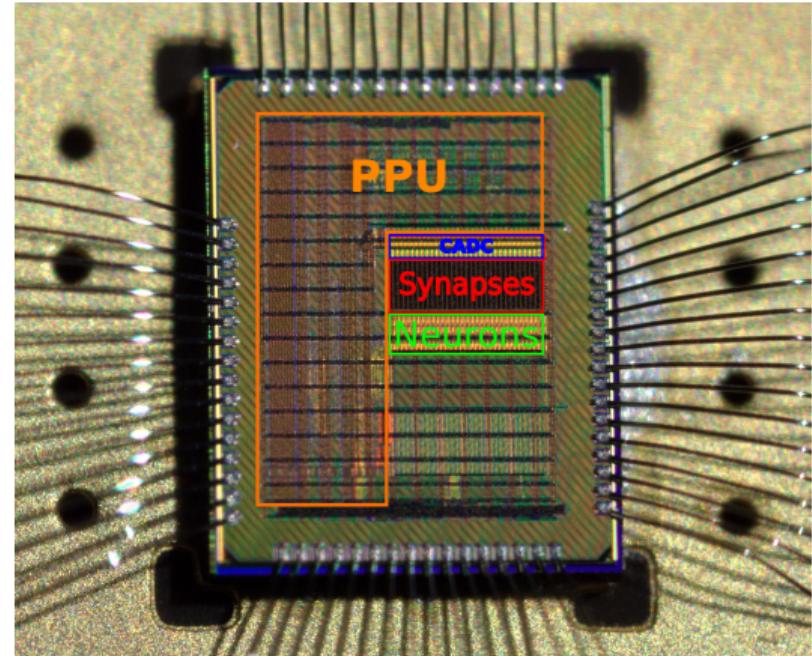


Figure 3: Photograph of the HICANN-DLS Chip, modified from Friedmann et al.

Plasticity Processing Unit

- ▶ “Common” von-Neumann architecture
- ▶ Machine Instructions
- ▶ Arithmetic Logic Unit (ALU)

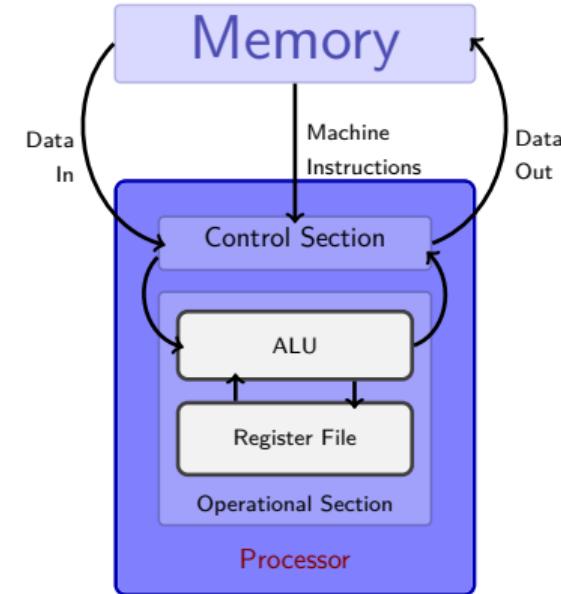


Figure 4: Schematic of General Purpose Processor with von-Neumann Architecture

Plasticity Processing Unit

- ▶ “Common” von-Neumann architecture
- ▶ Machine Instructions
- ▶ Arithmetic Logic Unit (ALU)
- ▶ $\text{latency}(\text{register}) \ll \text{latency}(\text{memory})$

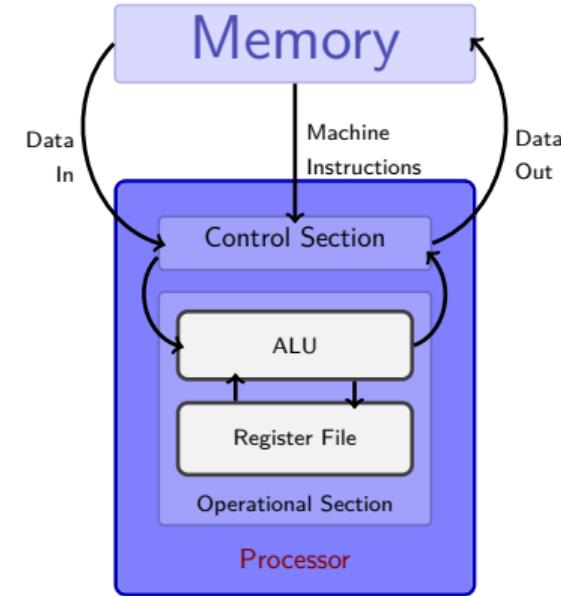


Figure 4: Schematic of General Purpose Processor with von-Neumann Architecture

PPU Architecture

- ▶ Based on POWER architecture
- ▶ Vector Extension
 - Parallelization
 - Performance
- ▶ Weight Updates
- ▶ Access to Synaptic Array

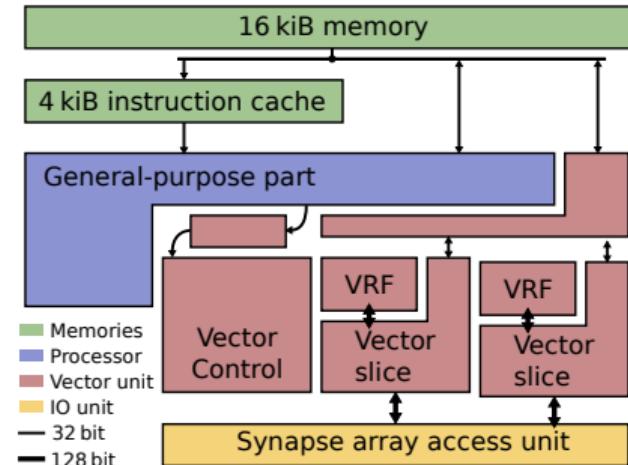


Figure 5: Structure of nux Architecture

Compiler Structure

- ▶ Program $\xrightarrow{\text{compile}}$ Executable
- ▶ Compiling Stages
- ▶ Compiler → Assembly
- ▶ Back-End is target-dependent

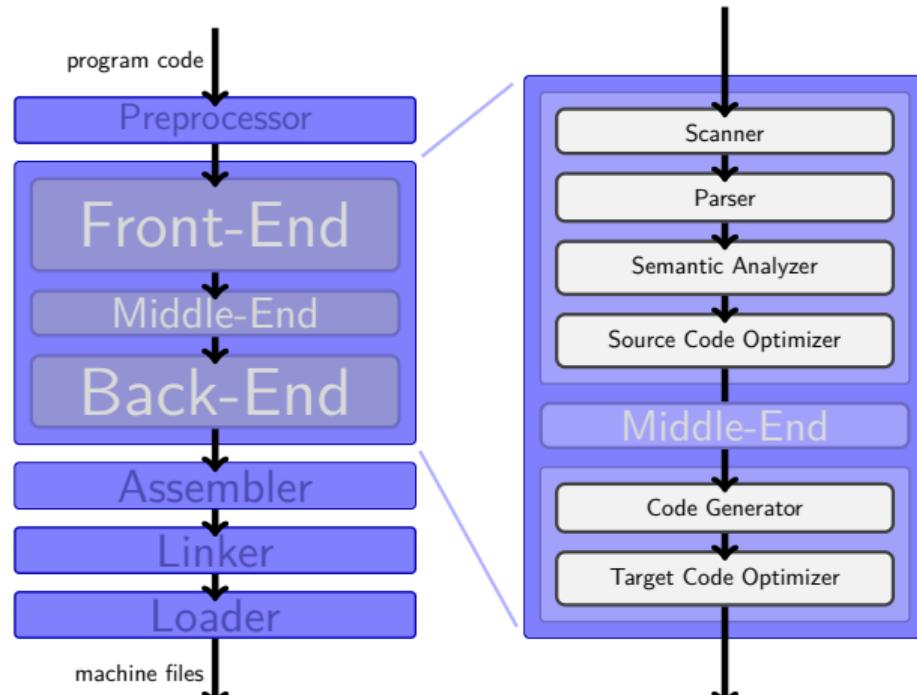


Figure 6: Structure of Compiling Process and Compiler

GCC for the PPU

- ▶ No “official” support for s2pp
- ▶ Currently using macros

Example from ppu_sweep.c

```
1 ...  
2 /* Definitions for vector registers */  
3 #define VR_TMP 0  
4 #define VR_CAUSAL 1  
5 #define VR_ACAUSAL 2  
6 #define VR_RST 3  
7 #define VR_NULL_C 4  
8  
9 ...
```

GCC for the PPU

- ▶ No “official” support for s2pp
- ▶ Currently using macros

Example of Macro Usage

```
1 ...  
2   fvx_splatb(VR_TMP, 1);  
3   fvx_addbm(VR_TMP_1, VR_A, VR_TMP);
```

GCC for the PPU

- ▶ No “official” support for s2pp
- ▶ Currently using macros
- ▶ Existing AltiVec extension

Example of Macro Usage

```
1 ...  
2 fxv_splatb(VR_TMP, 1);  
3 fxv_addbm(VR_TMP_1, VR_A, VR_TMP);
```

Example of AltiVec

```
1 ...  
2 tmp = vec_splat(1);  
3 tmp = vec_add(tmp, tmp);
```

GCC for the PPU

- ▶ No “official” support for s2pp
- ▶ Currently using macros
- ▶ Existing AltiVec extension
- ▶ Intrinsics → **use this**

Example of Macro Usage

```
1 ...  
2 fxv_splatb(VR_TMP, 1);  
3 fxv_addbm(VR_TMP_1, VR_A, VR_TMP);
```

Example of AltiVec

```
1 ...  
2 tmp = vec_splat(1);  
3 tmp = vec_add(tmp, tmp);
```

Main Work of this Thesis

- ▶ Add GCC support for s2pp
- ▶ Overcome missing documentation
- ▶ AltiVec inspiration
- ▶ Components:
 - Vector Register
 - `vector` attribute
 - New intrinsics
 - ...

Example of Macro Usage

```
1 ...  
2 fxv_splatb(VR_TMP, 1);  
3 fxv_addbm(VR_TMP_1, VR_A, VR_TMP);
```

Example of AltiVec

```
1 ...  
2 tmp = vec_splat(1);  
3 tmp = vec_add(tmp, tmp);
```

New Features

- ▶ -mcpu=nux target flag

Compiling Directive

```
1 powerpc-linux-eabi-ppc -mcpu=nux
```

Example File

```
1 #include<s2pp.h>
2
3 start(){
4     return;
5 }
```

New Features

- ▶ `-mcpu=nux` target flag
- ▶ `vector` attribute

Compiling Directive
1 `powerpc-linux-eabi-ppc -mcpu=nux`

Example File

```
1 #include<s2pp.h>
2
3 start(){
4     vector int8_t vec;
5     return;
6 }
```

New Features

- ▶ `-mcpu=nux` target flag
- ▶ `vector` attribute
- ▶ s2pp vector intrinsics

Compiling Directive
1 `powerpc-linux-eabi-ppc -mcpu=nux`

Example File
1 `#include<s2pp.h>`
2
3 `start(){`
4 `vector int8_t vec =fxv_splat(1);`
5 `vec = fxv_add(vec, vec);`
6 `return;`
7 `}`

New Features

- ▶ `-mcpu=nux` target flag
- ▶ `vector` attribute
- ▶ s2pp vector intrinsics
- ▶ Inline assembly

Compiling Directive
1 `powerpc-linux-eabi-ppc -mcpu=nux`

Example File
1 `#include<s2pp.h>`
2
3 `start(){`
4 `vector int8_t vec =fxv_splat(1);`
5 `vec = fxv_add(vec, vec);`
6 `asm ("fxvaddbm %0, %0, %0"
7 ":=kv" (vec):::);`
8 `return;`
9 `}`

New Features

- ▶ `-mcpu=nux` target flag
- ▶ `vector` attribute
- ▶ s2pp vector intrinsics
- ▶ Inline assembly
- ▶ Supports optimization

Compiling Directive

```
1 powerpc-linux-eabi-ppc -mcpu=nux -O1
```

Example File

```
1 #include<s2pp.h>
2
3 start(){
4     vector int8_t vec =fxv_splat(1);
5     vec = fxv_add(vec, vec);
6     asm ("fxvaddbm %0, %0, %0"
7          :"=kv" (vec)::);
8     return;
9 }
```

Comparison of unoptimized code (left) and **-O1** optimized code (right)

1 start:	1 start:
2 stwu %r1,-64(%r1)	2 stwu %r1,-48(%r1)
3 stw %r31,60(%r1)	3 li %r9,1
4 mr %r31,%r1	4 fxv splatb %f12,%r9
5 li %r9,1	5 li %r9,16
6 fxv splatb %f12,%r9	6 fxv stax %f12,%r1,%r9
7 li %r9,16	7 fxv lax %f12,%r1,%r9
8 fxv stax %f12,%r31,%r9	8 fxv lax %f11,%r1,%r9
9 li %r9,16	9 fxv addbfs %f12,%f12,%f11
10 fxv lax %f12,%r31,%r9	10 fxv stax %f12,%r1,%r9
11 li %r9,16	11 #APP
12 fxv lax %f11,%r31,%r9	12 # 277 "nux/main.c" 1
13 fxv addbfs %f12,%f12,%f11	13 fxv addbm %f12, %f12, %f12
14 li %r9,16	14 # 0 "" 2
15 fxv stax %f12,%r31,%r9	15 #NO_APP
16 #APP	16 fxv stax %f12,%r1,%r9
17 # 277 "nux/main.c" 1	17 addi %r1,%r1,48
18 fxv addbm %f12, %f12, %f12	18 blr
19 # 0 "" 2	
20 #NO_APP	
21 li %r9,16	
22 fxv stax %f12,%r31,%r9	
23 nop	
24 addi %r11,%r31,64	
25 lwz %r31,-4(%r11)	
26 mr %r1,%r11	
27 blr	

Outlook

- ▶ Further testing for bugs
- ▶ Existing applications with `linux`
- ▶ Extending test coverage
- ▶ Debugging support?
- ▶ **New tool for development!**

Comparison of unoptimized code (left) and `-O1` optimized code (right)

```
1  start:  
2  stwu %r1,-64(%r1)  
3  stw %r31,60(%r1)  
4  mr %r31,%r1  
5  li %r9,1  
6  fxv splatb %f12,%r9      1  start:  
7  li %r9,16                 2  stwu %r1,-48(%r1)  
8  fxv stax %f12,%r31,%r9   3  li %r9,1  
9  li %r9,16                 4  fxv splatb %f12,%r9  
10 fxv lax %f12,%r31,%r9    5  li %r9,16  
11 li %r9,16                 6  fxv stax %f12,%r1,%r9  
12 fxv lax %f11,%r31,%r9    7  fxv lax %f12,%r1,%r9  
13 fxv addbfs %f12,%f12,%f11 8  fxv lax %f11,%r1,%r9  
14 li %r9,16                 9  fxv addbfs %f12,%f12,%f11  
15 fxv stax %f12,%r31,%r9   10 fxv stax %f12,%r1,%r9  
16 #APP                      11 #APP  
17 # 277 "nux/main.c" 1      12 # 277 "nux/main.c" 1  
18 fxv addbbm %f12, %f12, %f12 13 fxv addbbm %f12, %f12, %f12  
19 # 0 "" 2                  14 # 0 "" 2  
20 #NO_APP                   15 #NO_APP  
21 li %r9,16                 16 fxv stax %f12,%r1,%r9  
22 fxv stax %f12,%r31,%r9   17 addi %r1,%r1,48  
23 nop                      18 blr  
24 addi %r11,%r31,64  
25 lwz %r31,-4(%r11)  
26 mr %r1,%r11  
27 blr
```

References

- S. Friedmann, J. Schemmel, A. Grübl, A. Hartel, M. Hock, and K. Meier, "Demonstrating hybrid learning in a flexible neuromorphic hardware system", (2016).
- T. Flik, *Mikroprozessortechnik und rechnerstrukturen*, ger, 7., neu bearb. Aufl. (Springer, Berlin ; Heidelberg [u.a.], 2005), XIV, 649 S.
- K. D. Cooper and L. Torczon, *Engineering a compiler*, eng, 2. ed., Previous ed.: 2004 (Elsevier, Amsterdam ; Heidelberg [u.a.], 2012), XXIII, 800 S.
- A. V. Aho, *Compiler, Prinzipien, techniken und werkzeuge*, ger, 2., aktualisierte Aufl., it Informatik, Index S. 1227-1253 (Pearson Studium, München [u.a.], 2008), XXXVI, 1253 S.
- S. Friedmann, *Nux manual*, (2016).
- *Gnu compiler collection internals manual*,
<https://gcc.gnu.org/onlinedocs/gccint/index.html>, Free Software Foundation Inc. (2017).

Whiteboard

7	6	5	4	3	2	1	0
		2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-6}
7	6	5	4	3	2	1	0
		2^5	2^4	2^3	2^2	2^1	2^0
7	6	5	4	3	2	1	0
-1	2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-5}	2^{-7}

Figure 7: Comparison of the Representation of Weights in Synapses and the Fractional Representation of Vector Components for Fixed-Point Saturation Arithmetic

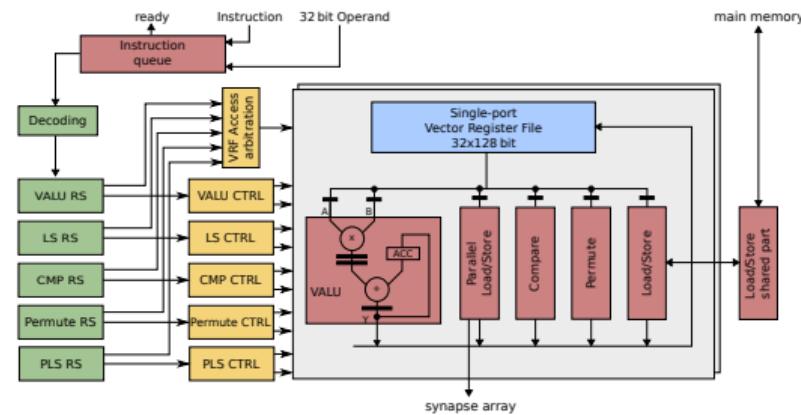
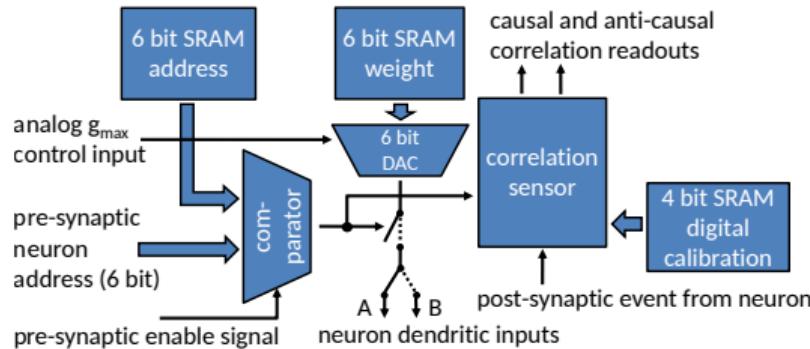
PPU is the processor which is part of HICANN-DLS and mainly responsible for plasticity.

nux refers to the architecture of the PPU, see figure.

s2pp describes the PPU's VE and is part of the nux architecture.

Appendix

Additional Figures



	31	23	15	7	0
mnemonic	operand	operand	operand		
addi	r1 register address	r2 register address	5 immediate operand		

Figure 9: Representation of Assembly Instruction addi as a Machine Instruction in Memory. The immediate value 5 is added to register r2 and the result written in r1.

Additional Figures

		0	7	15									31									63									127
QI		Quarter Integer																													
HI		Half Integer																													
SI														Single Integer																	
SF														Single Float																	

Figure 11: Vector structures are 128 bits wide and split into common word sizes.

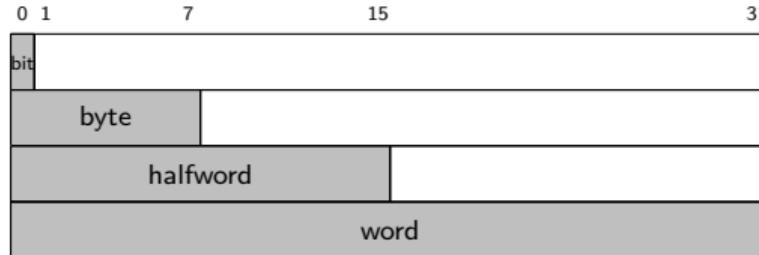


Figure 12: Illustration of Word Sizes for 32-bit Words